Elastic Load Balancer

Elastic Load Balancing distributes incoming application traffic across multiple EC2 instances, in multiple Availability Zones. This increases the fault tolerance of your applications. A load balancer accepts incoming traffic from clients and routes requests to its registered EC2 instances in one or more Availability Zones. The load balancer also monitors the health of its registered instances and ensures that it routes traffic only to healthy instances. When the load balancer detects an unhealthy instance, it stops routing traffic to that instance, and then resumes routing traffic to that instance when it detects that the instance is healthy again.

Elastic Load Balancing supports two types of load balancers:

Application Load Balancers and Classic Load Balancers. There is a key difference between the way you configure these load balancers. With a Classic Load Balancer, you register instances with the load balancer. With an Application Load Balancer, you register the instances as targets in a target group, and route traffic to a target group.

When you create a load balancer, you must choose whether to make it an internal load balancer or an Internet-facing load balancer. Note that when you create a Classic Load Balancer in EC2Classic, it must be an Internet-facing load balancer. The nodes of an Internet-facing load balancer have public IP addresses. The DNS name of an Internet facing load balancer is publicly resolvable to the public IP addresses of the nodes. Therefore, Internet facing load balancers can route requests from clients over the Internet.

The nodes of an internal load balancer have only private IP addresses. The DNS name of an internal load balancer is publicly resolvable to the private IP addresses of the nodes. Therefore, internal load balancers can only route requests from clients with access to the VPC for the load balancer.

Note: Both Internet-facing and internal load balancers route requests to your instances using private IP addresses. Therefore, your instances do not need public IP addresses to receive requests from an internal or an Internet-facing load balancer.

If your application has multiple tiers, for example web servers that must be connected to the Internet and database servers that are only connected to the web servers, you can design an architecture that uses both internal and Internet-facing load balancers. Create an Internet-facing load balancer and register the web servers with it. Create an internal load balancer and register the database servers with it.

The web servers receive requests from the Internet-facing load balancer and send requests for the database servers to the internal load balancer. The database servers receive requests from the internal load balancer.

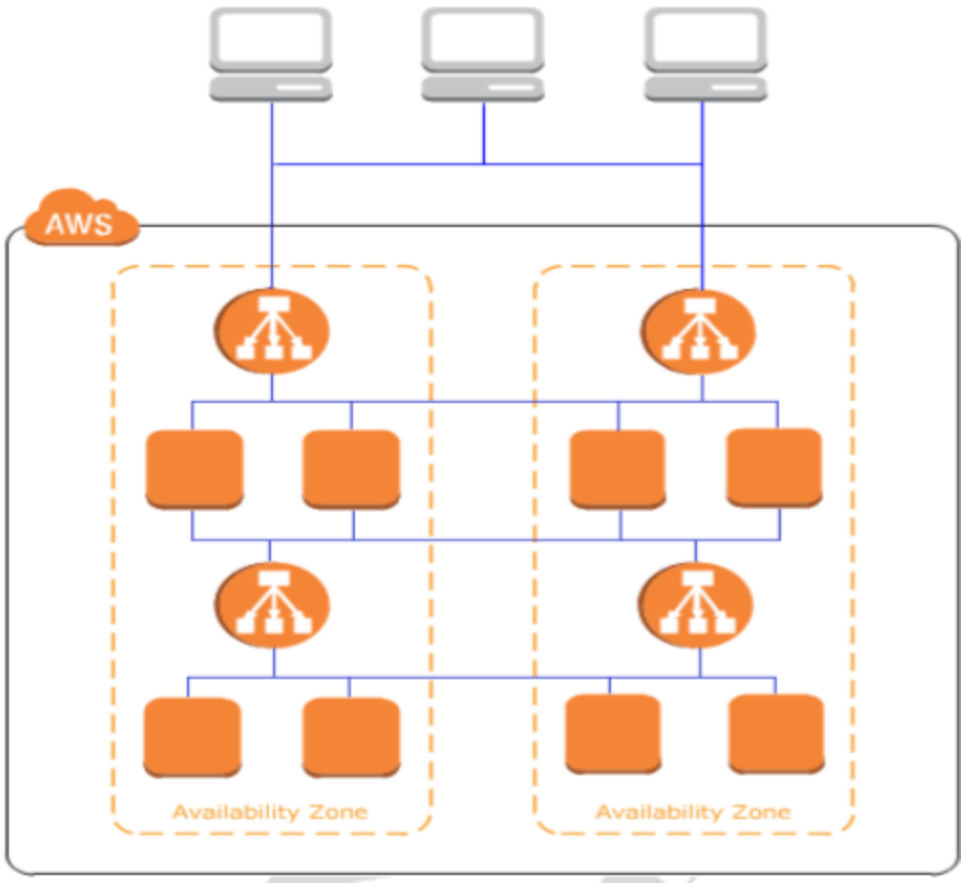

Creating Elastic Load Balancer:

Prerequisites:

• Choose any two Availability Zones you will use for your EC2 instances. Verify that your virtual private cloud (VPC) has at least one public subnet in each of these Availability Zones.

• Launch at least one EC2 instance in each Availability Zone.

• Ensure that the security group for your EC2 instances allows HTTP access on port 80. To test the web server, copy the DNS name of the instance and verify whether browser displays the default page of the web server or not.

The below screenshot shows the instance created named elbtestproj-web1, it has a public IP address but this IP is dynamic and changes after every reboot of the instance. We need to assign an Elastic IP to this instance which is static and does not change.

Elastic IP:

An Elastic IP address is a static IPv4 address designed for dynamic cloud computing. An Elastic IP address is associated with your AWS account. With an Elastic IP address, you can mask the failure of an instance or software by rapidly remapping the address to another instance in your account.

Assigning Elastic IP:

Navigate to the Elastic IPs tab on the EC2 Dashboard and click on allocate new address.

You will get a message ‘New address request succeeded’

Click on actions and select Associate address

Provide the instance id and click on associate. Then your instance will be allocated with this elastic IP

Login to the instance by using public IP

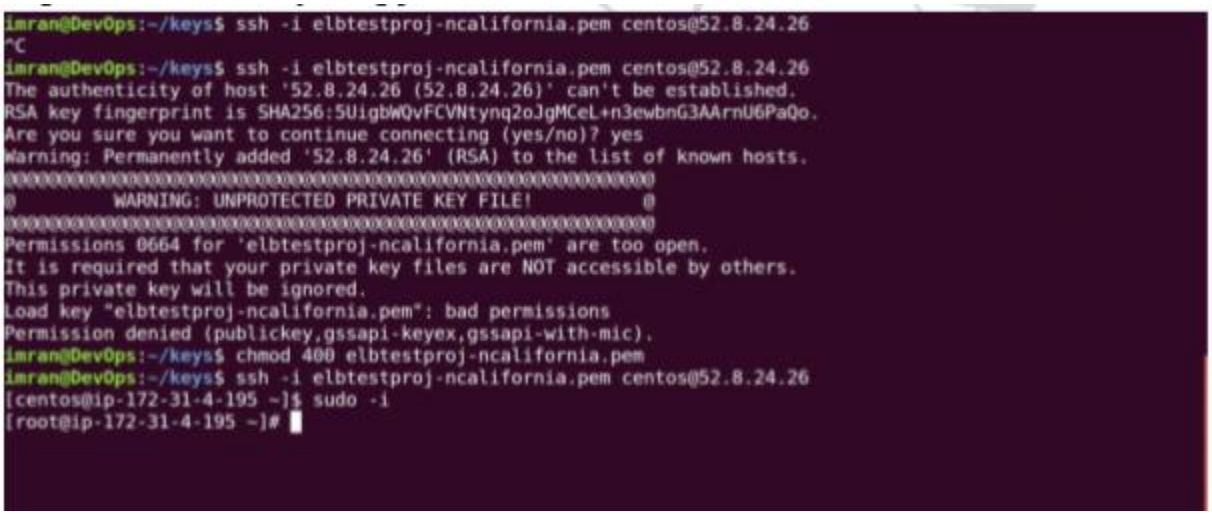

Create centos/6 EC2 instance.

Before that just create a new key pair and new security group and assign those to new EC2 when creating.

Install apache service by using below command

yum install httpd

Start apache service

service httpd start

Enable apache service

chkconfig httpd on

Stop & Disable firewall

service iptables stop

chkconfig iptables off

Create a test webpage for apache using html.

cd /var/www/html/ vi index.html

/var/www/html location is the default location to serve the web pages. Here we have an config file as well

Vim /etc/httpd/conf/httpd.conf -- all config about our apache server we have books to know this

We are going to have log files reg this apache and if we want to change any configuration we can change here like we can document root dir called /var/www/html location also.

Cd /var/log/httpd -- we get starting point to get trouble shoot when we have an error to start service

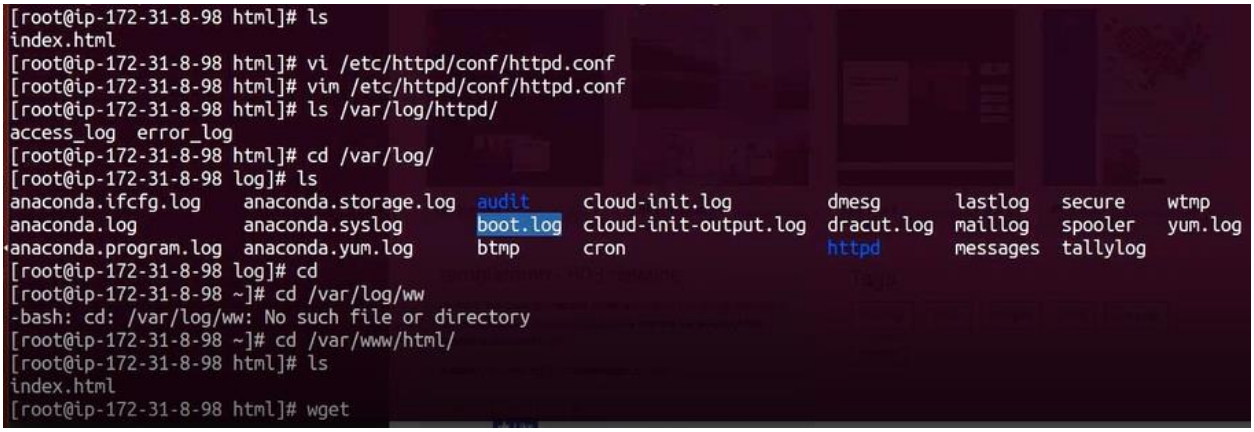

If we see any error we have to go to log folder and check for log file then only we can able to resolve the issue.

ls

Cd /var/www/html

<html>

<head>

<body>

<h1>This is a test webpage.</h1>

</body>

</head>

</html>

Or

Go to www.tooplate.com and download one static website and put into our server.

cd /var/www/html

sudo yum install wget unzip –y

wget –O website.zip <url from this site>

```
[root@ip-172-31-8-98 html]# wget -O website.zip http://www.templatemo.com/download/templatemo_503_newline
```

unzip website.zip

mv <dir name>/* .

```
[root@ip-172-31-8-98 html]# ls
templatemo_503_newline  website.zip
[root@ip-172-31-8-98 html]# pwd
/var/www/html
[root@ip-172-31-8-98 html]# ls
templatemo_503_newline  website.zip
[root@ip-172-31-8-98 html]# mv templatemo_503_newline/* .
[root@ip-172-31-8-98 html]# ls
css  fonts  img  index.html  js  templatemo_503_newline  videos  website.zip
[root@ip-172-31-8-98 html]#
```

```
[root@ip-172-31-8-98 html]# rm templatemo_503_newline/ -r
rm: remove directory `templatemo_503_newline'? y
[root@ip-172-31-8-98 html]# rm -rf website.zip
[root@ip-172-31-8-98 html]# ls
css  fonts  img  index.html  js  videos
[root@ip-172-31-8-98 html]#
```

Now just restart our web server.

sudo service apache2 stop

sudo service apache2 start

sudo service apache2 status

Test the webpage Enter the ec2 inst public IP in browser.

http://52.8.24.26:8080

We are able to see our static website into our server and it is better than a static text on our website.

AWS AMI:

An Amazon Machine Image (AMI) is a special type of virtual appliance that is used to create a virtual machine within the Amazon Elastic Compute Cloud ("EC2"). It serves as the basic unit of deployment for services delivered using EC2.

AMI Creation:

Create AMI of the instance which we will use to spin web02 instance. Web02 instance is exactly similar to web01, so instead to creating new instance from scratch and setting up apache, we can create an AMI (image) of web01 instance and can spin as many as web instances we want. Select the instance in which we have to create an image. Click on actions, select image and click on create image.

- You will get a create image dialog box as shown below. Give proper name and description for image and click on create image
- It takes few minutes to create an image which you can see on AMIs navigation pane

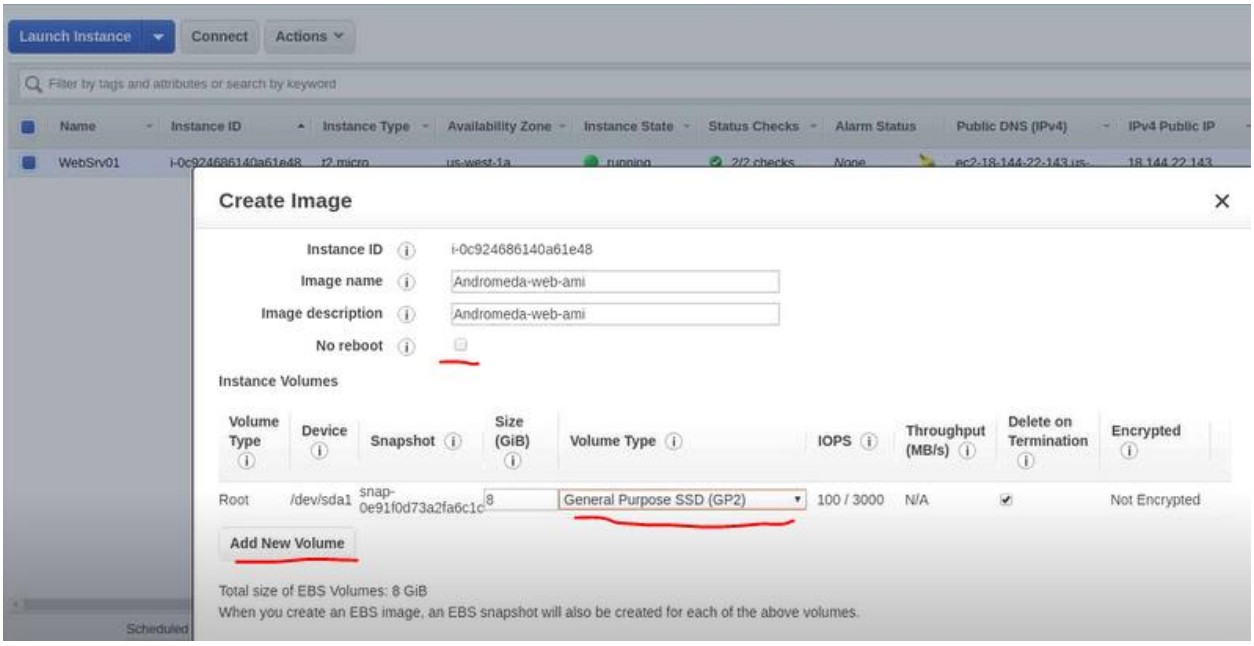

Create web02 instance from elbtestproj-web-ami.

Click on Launch instance --> My AMI --> Select your AMI --> Follow the wizard and create the instance similar to web01.

Tag Name: elbtestproj-web02 Security: Select existing security group --> elbtestproj-SG Select an existing key pair --> <same key used for web01> Assign Elastic IP

Test the webpage Enter the ec2 instance public IP in browser.

http://54.215.191.250//

It is just a snapshot, but snapshot is all about volume where as AMI is all about complete EC2 instance.

Load Balancer setup :

First we have to create target group for instances and then load balancer has to create.

- Click on Target Group on the left side of the navigation pane, click on create target group.
- Provide name for target group and click on create
- Once the Target group is created click on the targets tab, Edit and Select web01 & web02 instances. Add to registered and click on Save.
- Then the instances will be added to the Target Group.

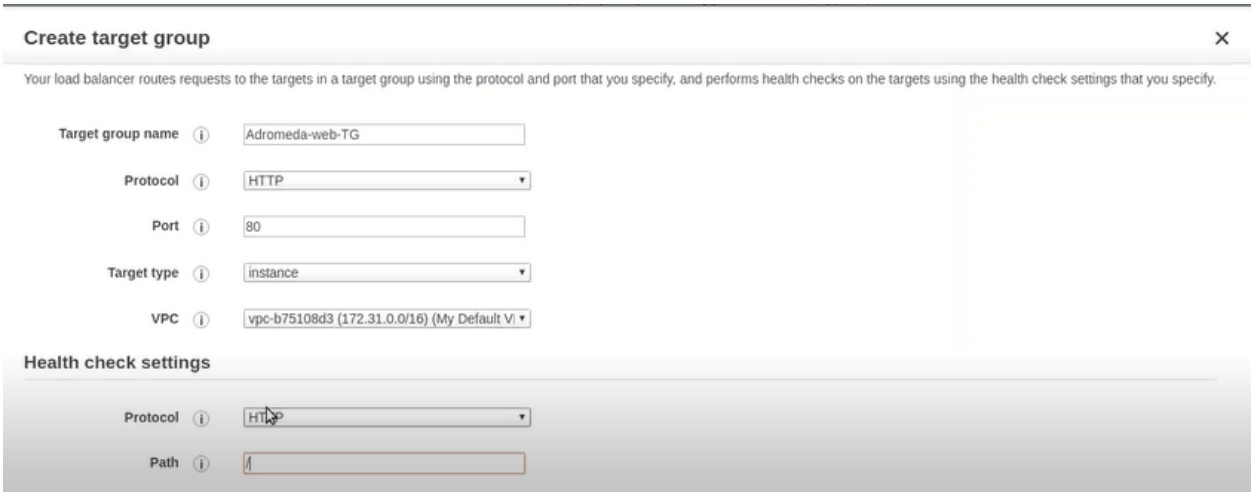

Now Target group is created and we have to add instances under target group here

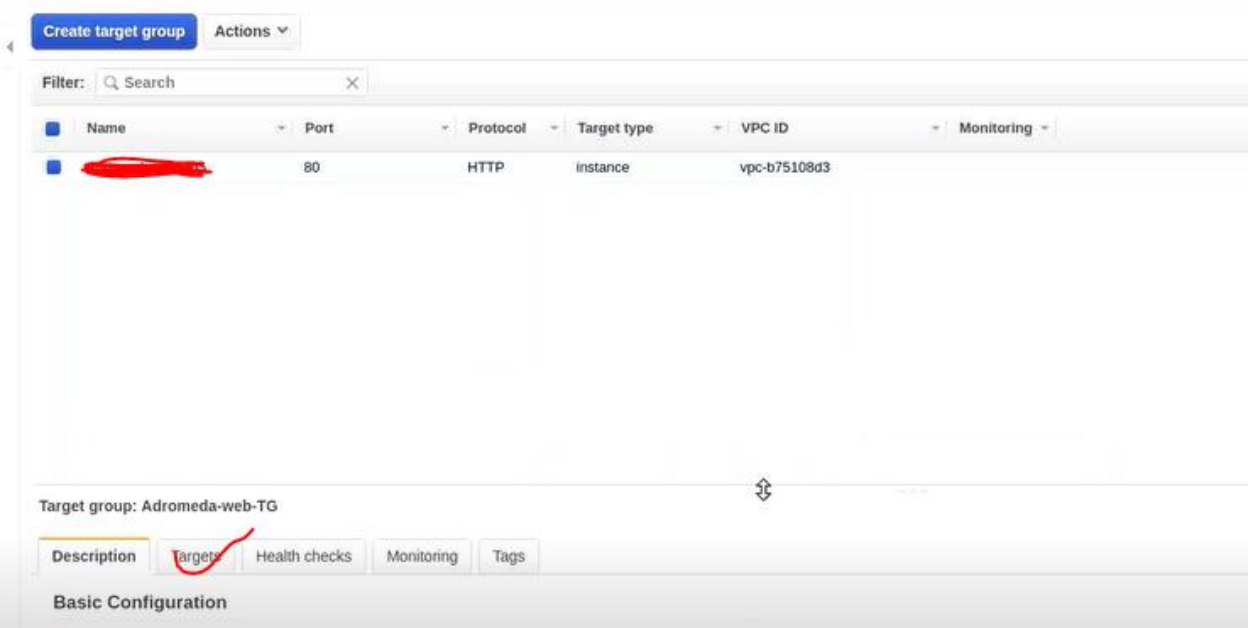

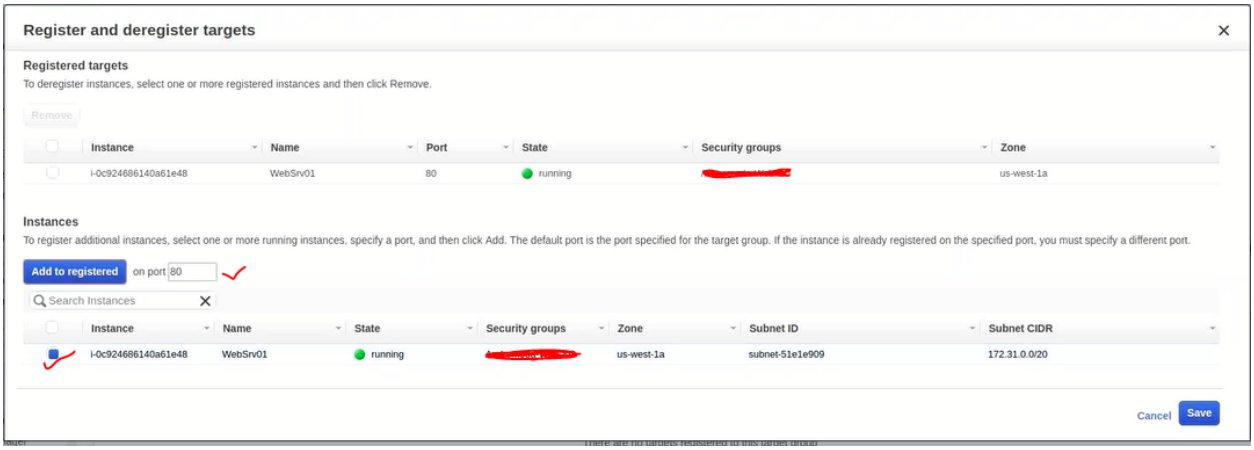

Added one instance and one mre we have to create from AMI and need to add here.

Creating Load Balancer:

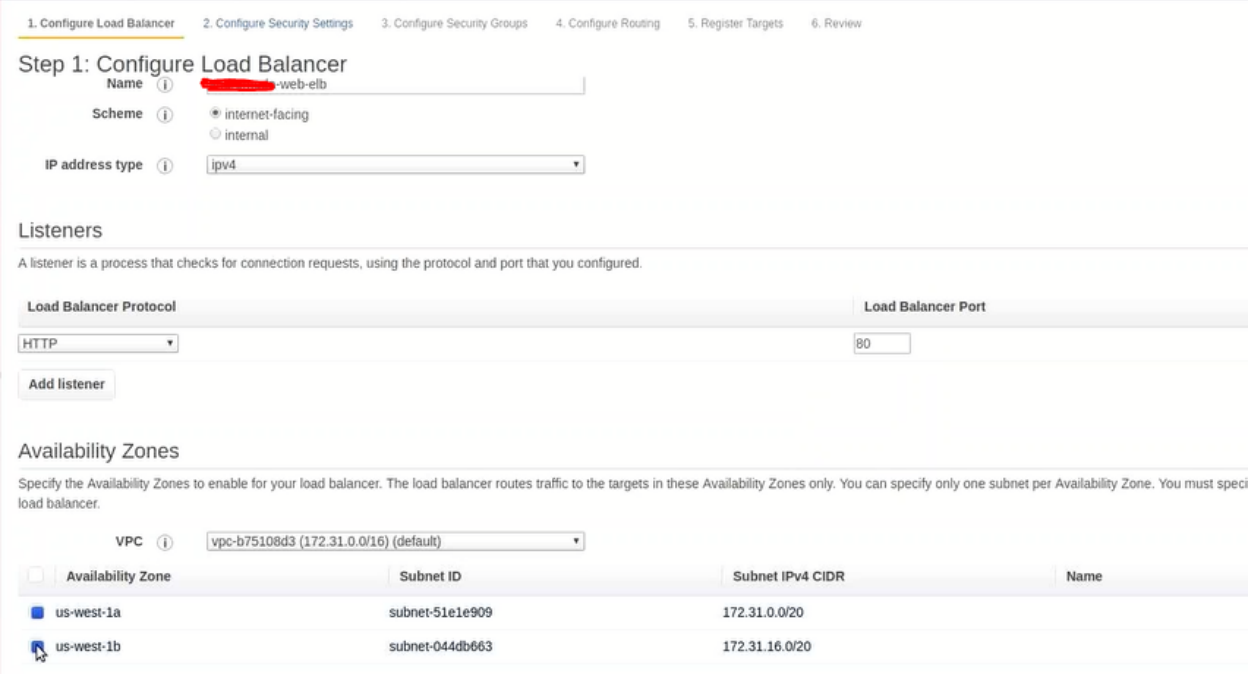

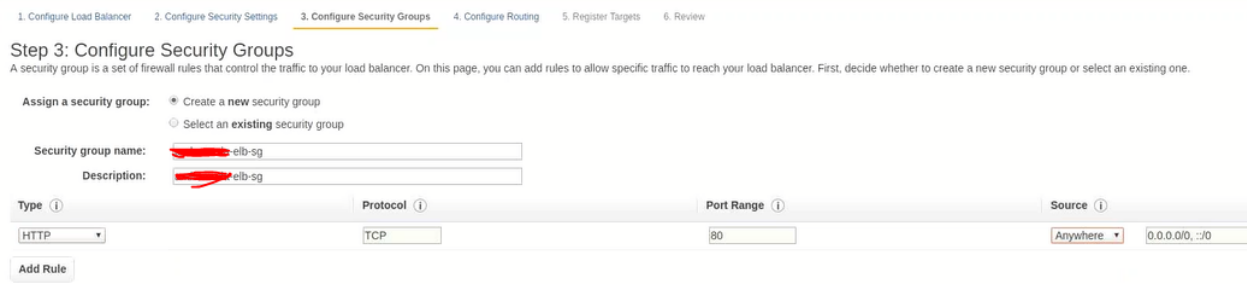

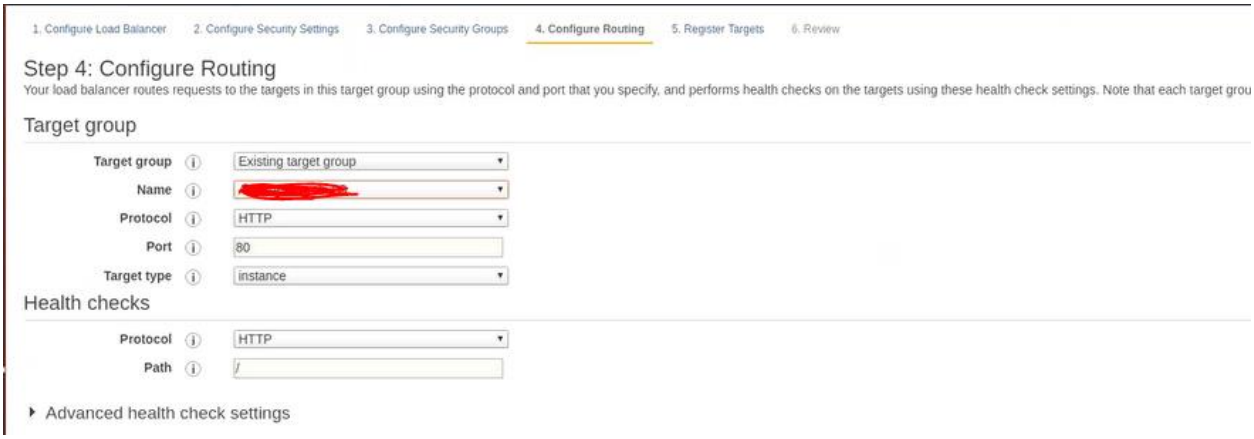

And Register and create. Done

Here am going to create one more instance from AMI and that is going to be add into this target group and accessing from this load balancer.

Her we can Health check tab under target group reg health check.

Click on Load Balancers in Left Pane and select Create Load Balancer, you will get two types of load balancers, Application load balancer and Classic load balancer. Select Application Load Balancer and click on Continue.

Configure the details of the load balancer. Provide the name for load balancer, the name of your Application Load Balancer must be unique within your set of Application Load Balancers. For one region, it can have a maximum of 32 characters and contain only alphanumeric characters and hyphens, and must not begin or end with a hyphen. Click on next configure security group.

For Availability Zones, select the VPC that you used for your EC2 instances. For each of the two Availability Zones that contain your EC2 instances, select the Availability Zone and then select the public subnet for that Availability Zone. Add atleast two Availability Zones to increase the availability of the load balancer. Click on next

Configure the details of security group to improve the load balancer's security.

Create a new security group for load balancer which accepts HTTP traffic on port 80.

Configure the target group. The default rule for your listener routes requests to the registered targets in this target group. The load balancer checks the health of targets in this target group using the health check settings defined for the target group.

Register targets with the Target Groups

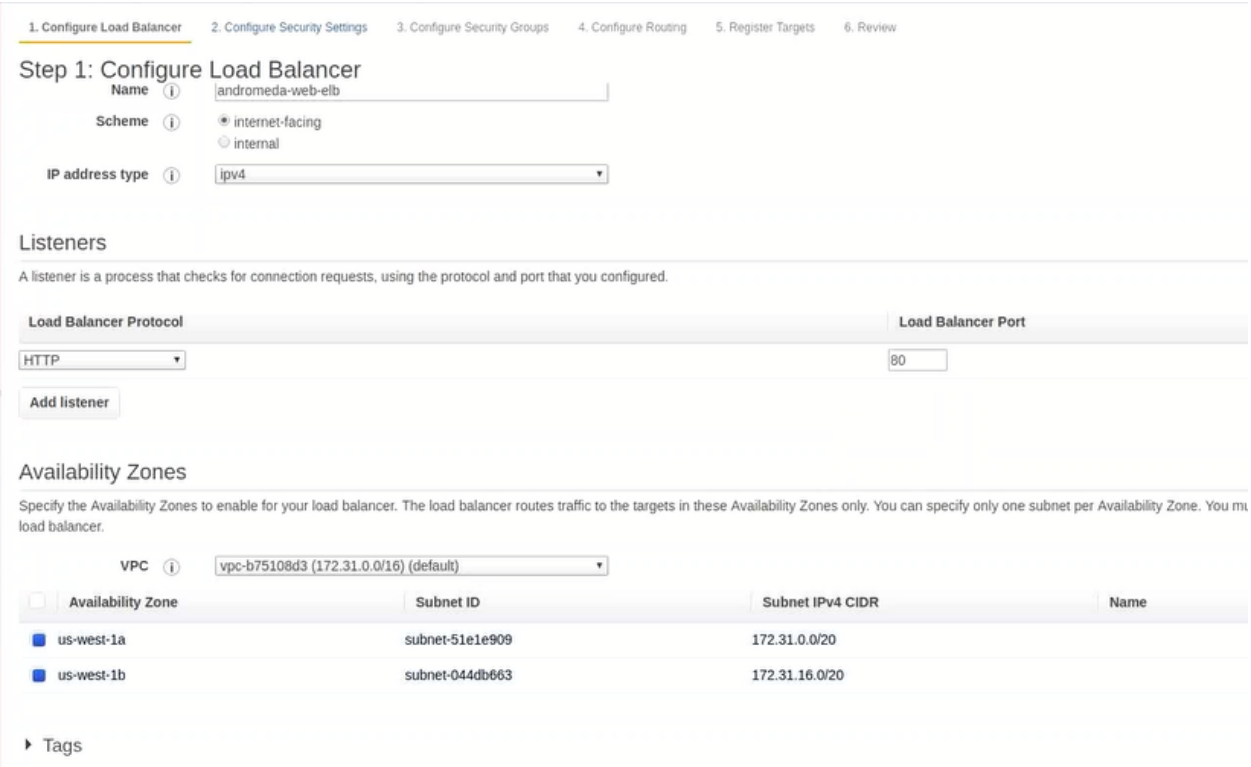

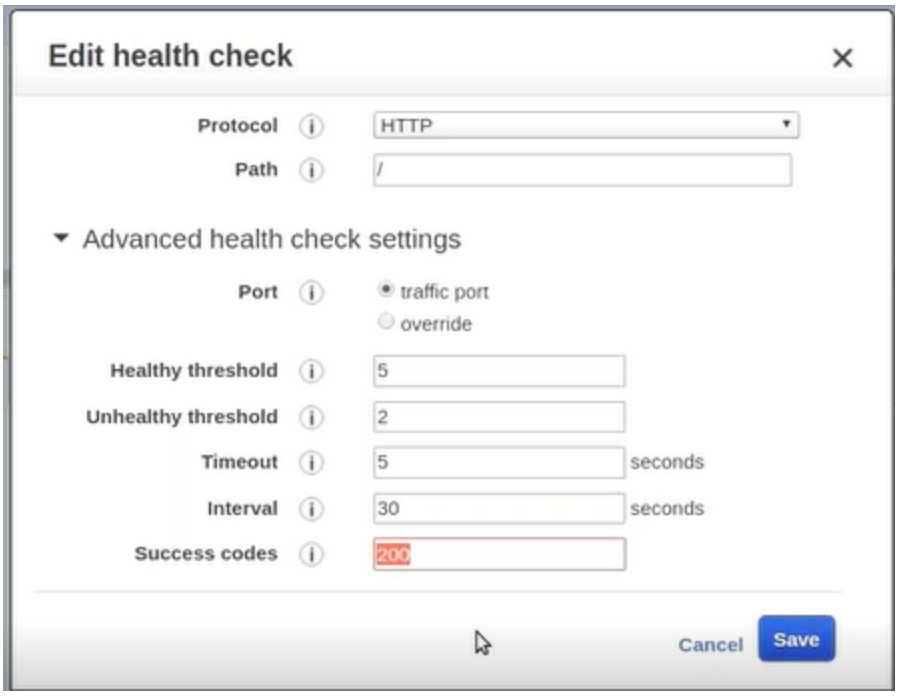

This is all about Health check about target group instances.